

**CSE202**  
**Design and Analysis of Algorithms**

***Week 10 — Balance against Worst-Case***

# Data-Structures for Ordered Data

Priority Queues: insert, findmax, deletemax

Ordered Search Trees: insert, find, delete, selectbyrank, floor, ceiling, countbetween,...

Sorting first is not an option

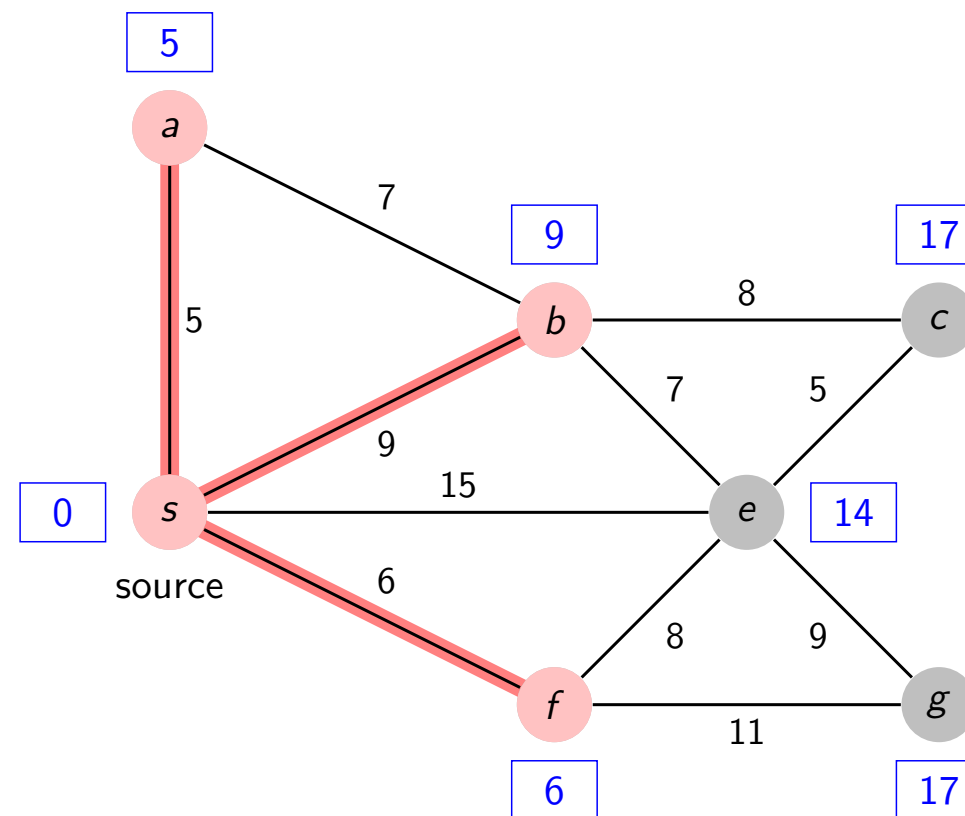
Def. All leaves in the same one or two levels

Balanced trees allow for all these operations in worst-case time  $O(\log n)$ .

$n$	$\log_2 n$
$10^6$	$\approx 20$
$10^9$	$\approx 30$
$10^{12}$	$\approx 40$

# **I. Priority Queues & Heap-ordered Trees**

# Recall Dijkstra's Algorithm (CSE103)



**while** PQ not empty:

remove first edge  $((u,v), d(s,u))$  from PQ

**if** v not in the tree

add v to the tree

**for all** neighbours w of v

insert  $((v,w), d(s,u)+d(v,w))$  in PQ

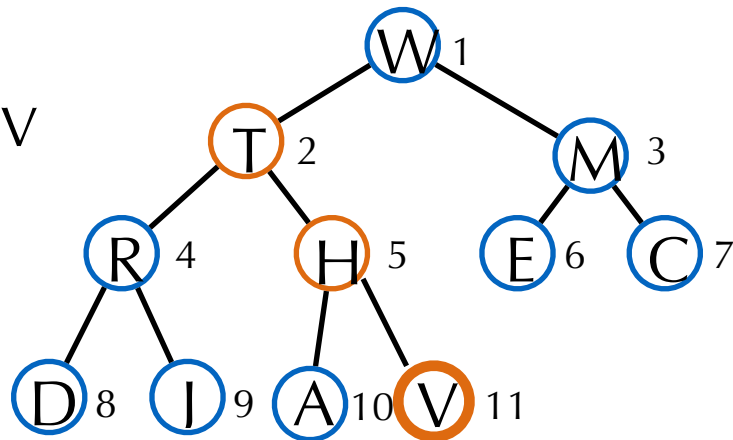
Complexity  
depends on good  
priority queues



# Basic Operations

## Insert & fixup

insert V



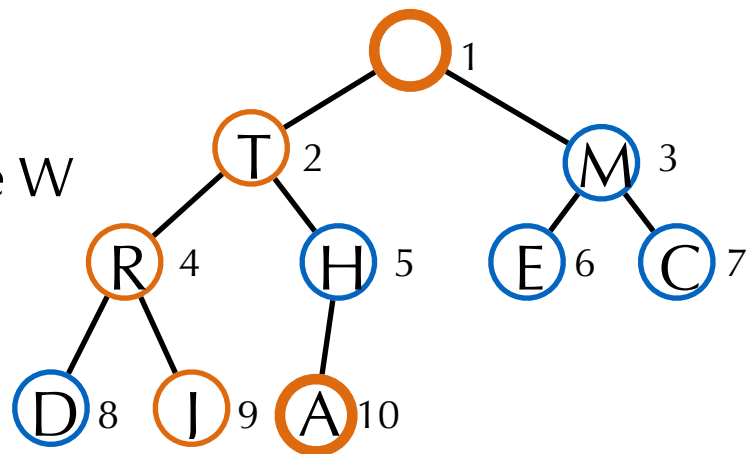
$\leq \log_2 n$  comparisons

```
def insert(self, key):  
    self.size += 1  
    self.PQ[self.size] = key  
    self.fixup(self.size)
```

```
def fixup(self, ind):  
    if ind == 1: return  
    parent = ind // 2  
    if self.PQ[parent] > self.PQ[ind]: return  
    self.exch(parent, ind)  
    self.fixup(parent)
```

## Deletemax & fixdown

delete W

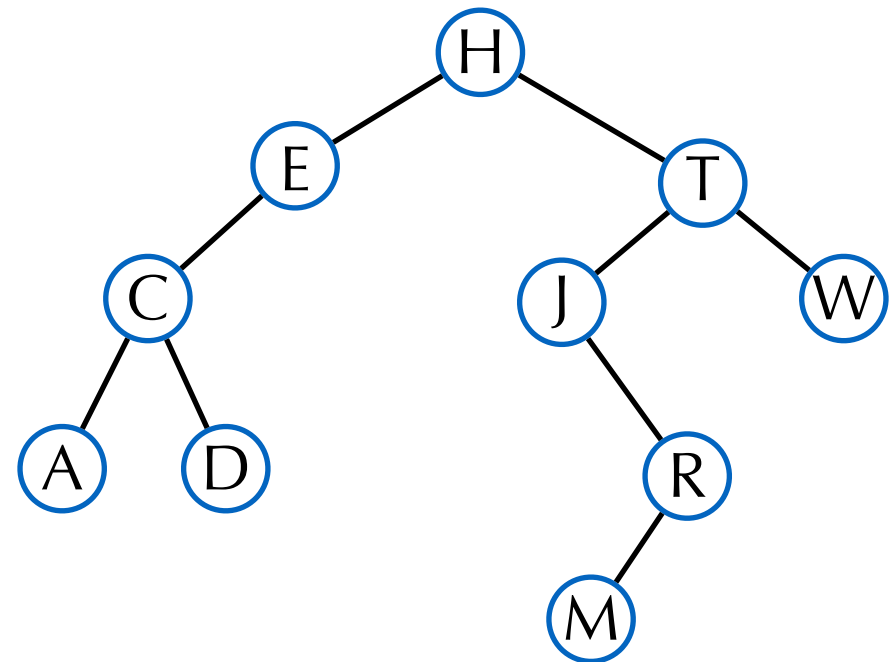


$\leq 2 \log_2 n$  comparisons

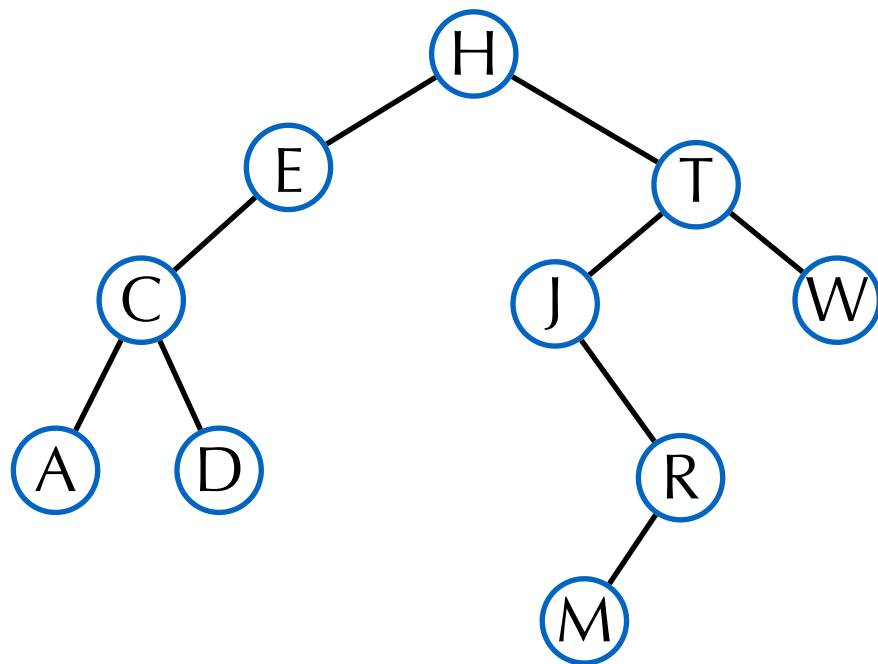
```
def deletemax(self):  
    self.PQ[1] = self.PQ[self.size]  
    self.size -= 1  
    self.fixdown(1)
```

```
def fixdown(self, ind):  
    child = 2 * ind  
    if child > self.size: return  
    if child < self.size and \  
        self.PQ[child + 1] > self.PQ[child]:  
        child += 1  
    if self.PQ[ind] < self.PQ[child]:  
        self.exch(ind, child)  
        self.fixdown(child)
```

## II. Binary Search Trees



# Recall Definition (CSE101 & 102)



Smaller elements to the left,  
larger elements to the right

```
class Node:
```

```
    def __init__(self, key, left=None, right=None):  
        self.key = key  
        self.left = left  
        self.right = right
```

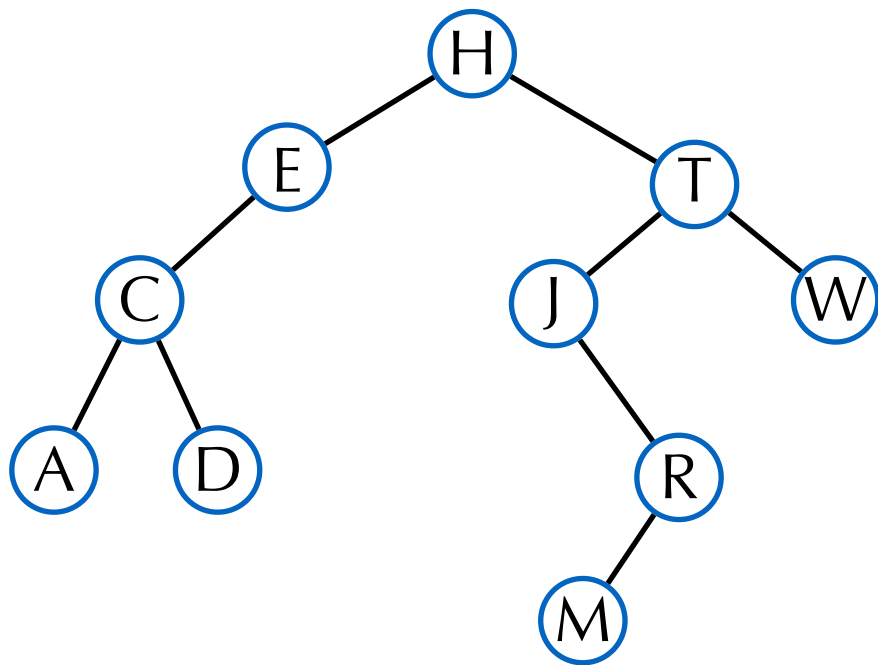
```
class BST:
```

```
    def __init__(self):  
        self.root = None  
  
    def find(self, key):  
        return self._find(self.root, key)  
  
    def insert(self, key):  
        self.root = self._insert(self.root, key)  
  
    def delete(self, key):  
        self.root = self._delete(self.root, key)
```



# Find/Insert

```
def _find(self, node, key):  
    if node is None: return False  
    if node.key > key: return self._find(node.left, key)  
    if node.key < key: return self._find(node.right, key)  
    return True
```



```
def _insert(self, node, key):  
    if node is None: return Node(key)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    return node
```

Delete slightly more complicated (CSE102)

**Worst-case:** search in  $O(n)$  comparisons for a BST built from  $n$  keys.

# Average-Case Analysis

Internal path length:

$P_n :=$  sum depths of all nodes

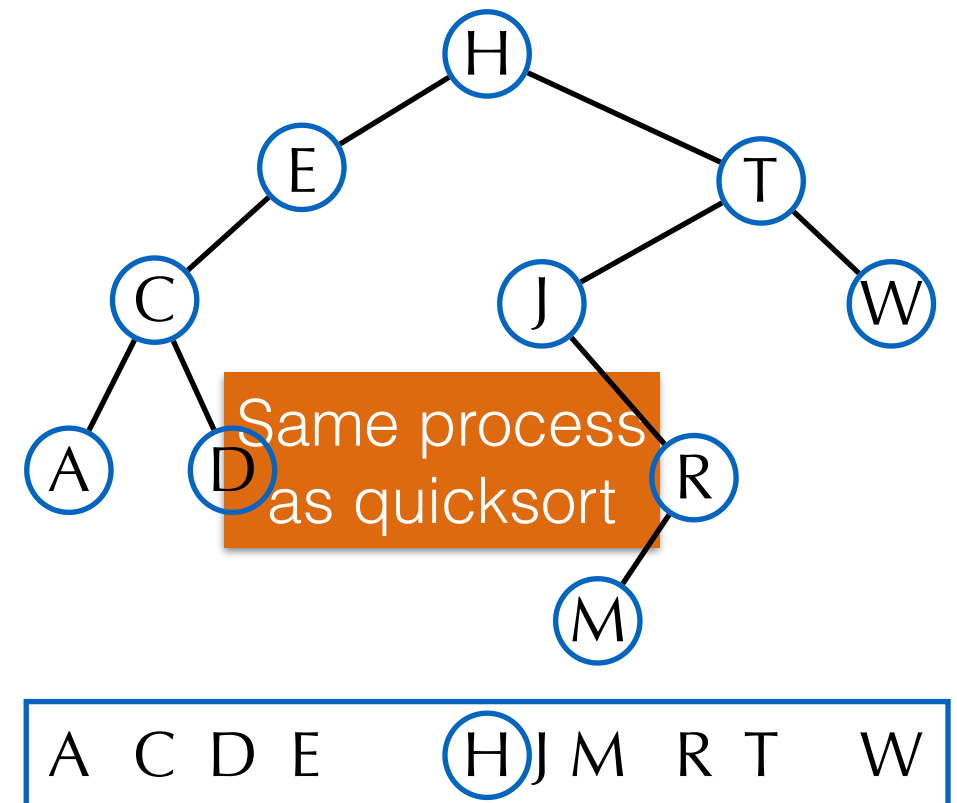
$P_n/n + 1$  : average successful search

$P_n/n + 3$  : average unsuccessful search  
(= insert)

Blackboard  
proof

**Prop.** In a BST built from  $n$  random keys, the average number of comparisons for a search is

$$2 \log n + O(1) \approx 1.39 \log_2 n$$

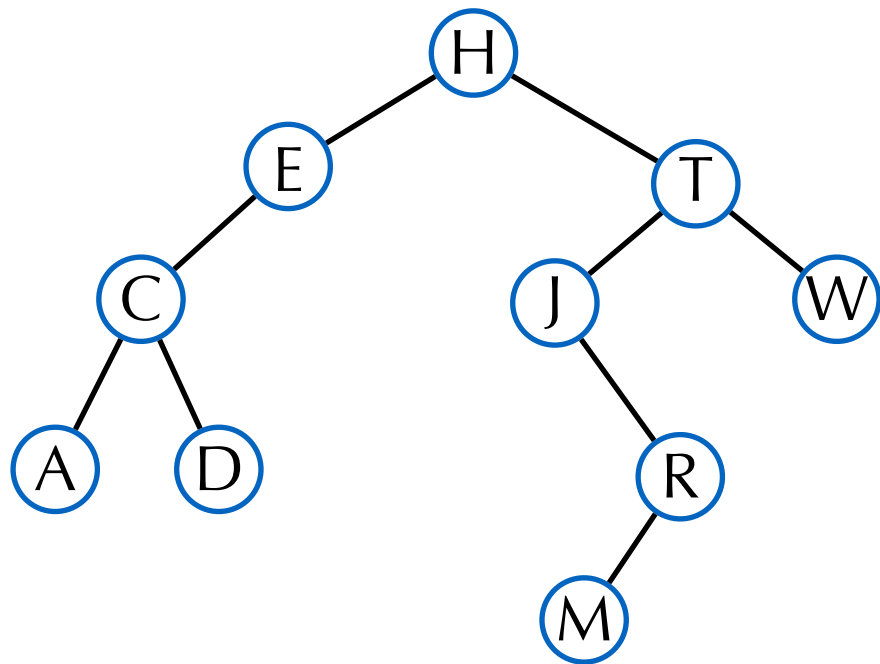


$$P_0 = P_1 = 0$$

$$\mathbb{E}P_n = n - 1 + \sum_{i=1}^n \frac{\mathbb{E}P_{i-1} + \mathbb{E}P_{n-i}}{n}$$

Same recurrence as in  
the analysis of quicksort.

# Select



min, max, floor, ceiling: easy

median, select:

change nodes into  
key, left, right, **size**

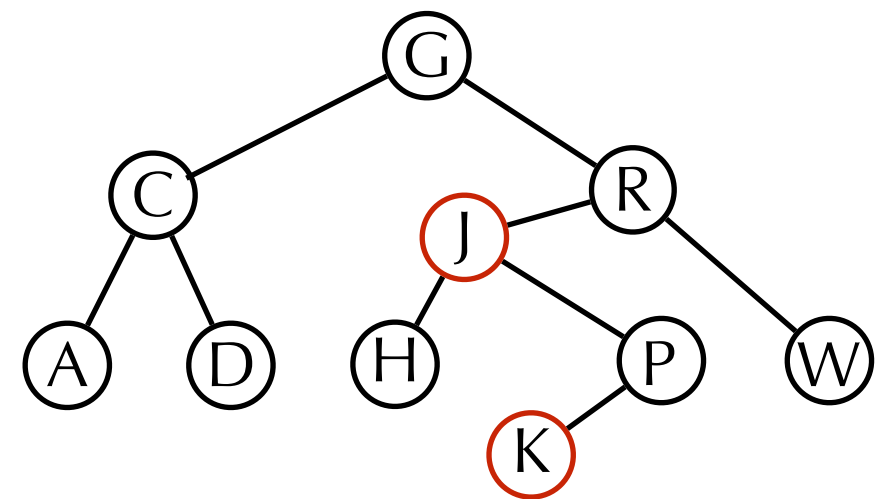
floor: largest key  
smaller than input

```
def _insert(self, node, key):  
    if node is None: return Node(key)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    node.size = 1 + size(node.left) + size(node.right)  
    return node
```

All these operations have cost bounded by the height,  
which is logarithmic **on average**.

Generalizes to  
higher dimensions  
(quadrees).

### III. Red-Black BST

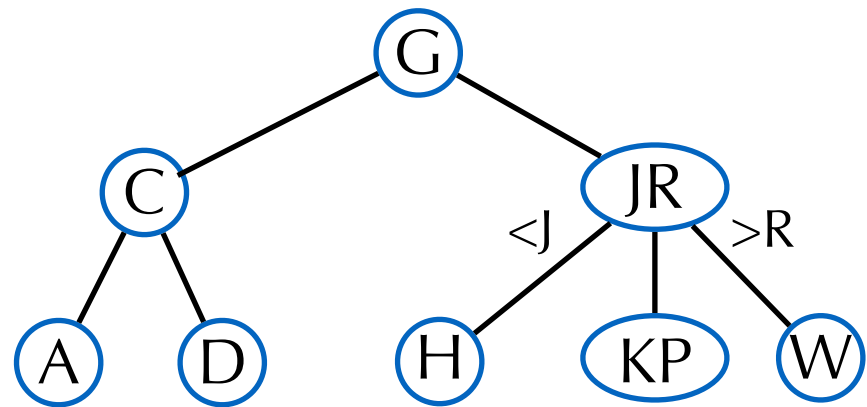


# WarmUp: 2-3 Search Trees

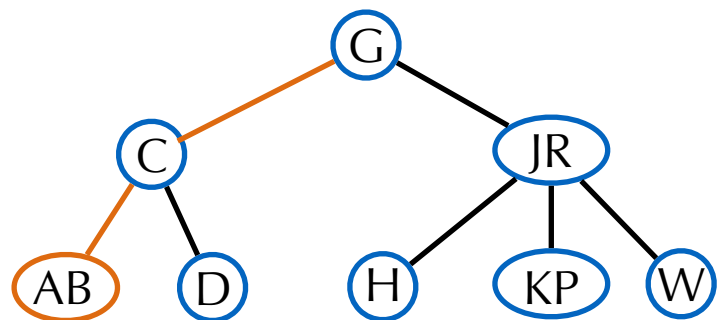
All leaves at the same level

Find: same as BST

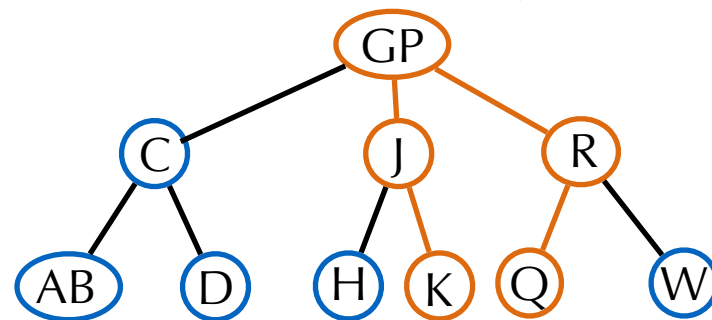
Insert maintaining perfect balance:  
search, insert at bottom  
and propagate upwards



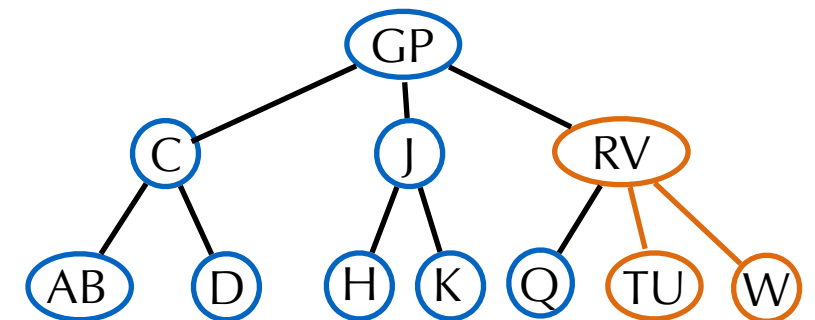
Insert B



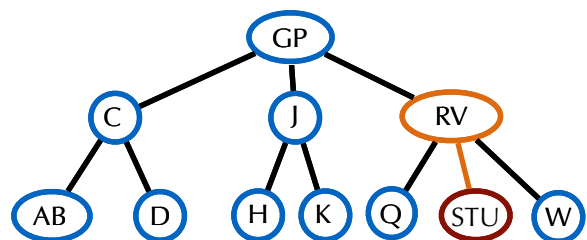
Insert Q



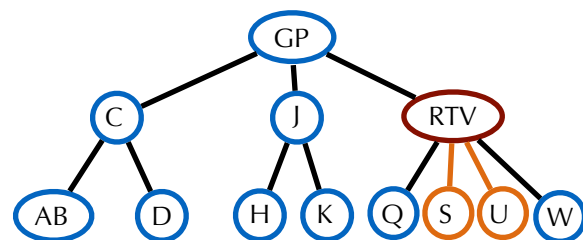
Insert V,U,T



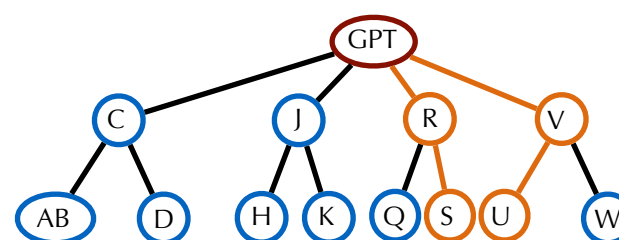
Insert S (1/4)



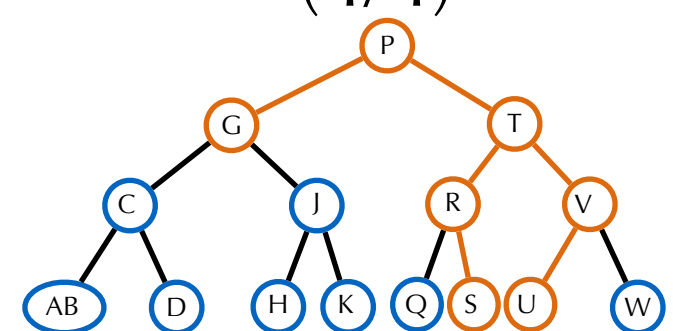
(2/4)



(3/4)



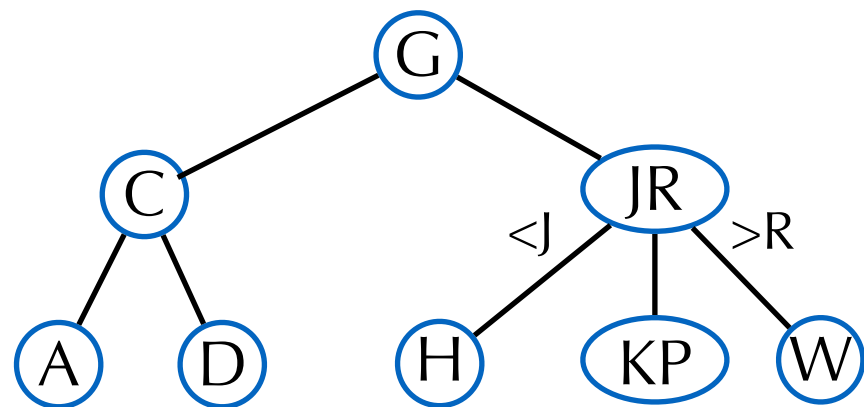
(4/4)



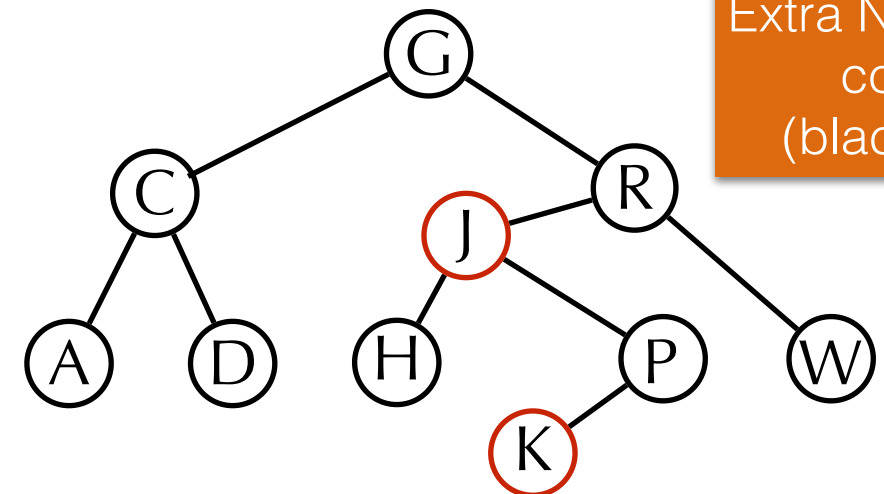
has to disappear

Worst-case number of nodes visited for find/insert:  $\log_2 n$

# (Left-Leaning) Red-Black Trees



stored as a  
coloured BST



Extra Node Info:  
colour  
(black root)

## Properties:

1. red nodes are left children;
2. red nodes have black children;
3. every path from the root to a leaf has the same number of black nodes.

Black  
balance

Red-black trees with these properties are in 1-to-1 correspondance with 2-3 trees.

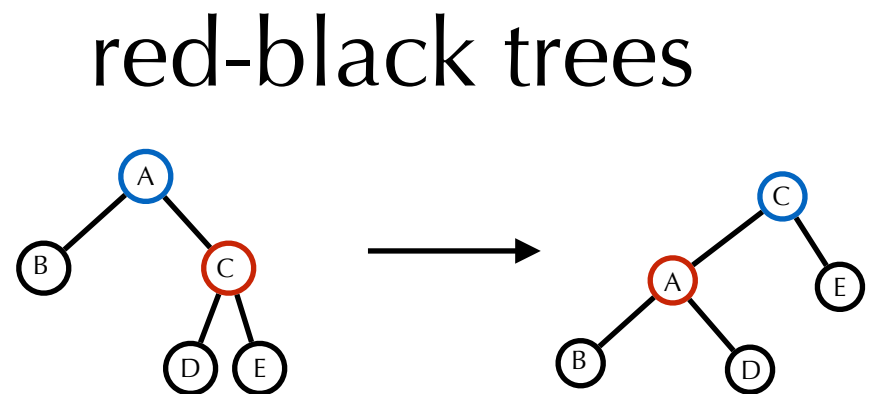
# Insertion

Insert maintaining  
order & black balance:  
search, insert **red** node at  
bottom and propagate upwards

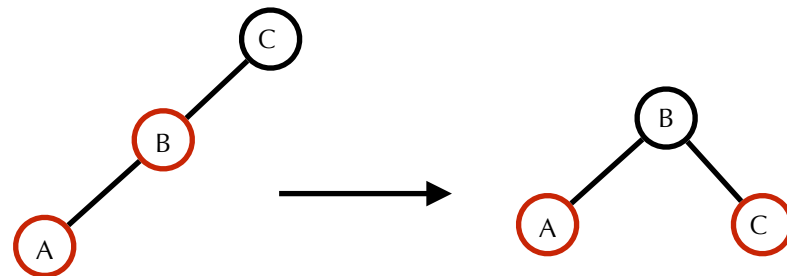
```
def _insert(self, node, key):  
    if node is None: return Node(key, red=True)  
    if node.key > key:  
        node.left = self._insert(node.left, key)  
    elif node.key < key:  
        node.right = self._insert(node.right, key)  
    if isRed(node.right) and not isRed(node.left): node = rotateleft(node)  
    if isRed(left.red) and isRed(node.left.left): node = rotateright(node)  
    if isRed(node.left) and isRed(node.right): flipcolors(node)  
    node.size = 1+size(node.left)+size(node.right)  
    return node
```

## Local fixes

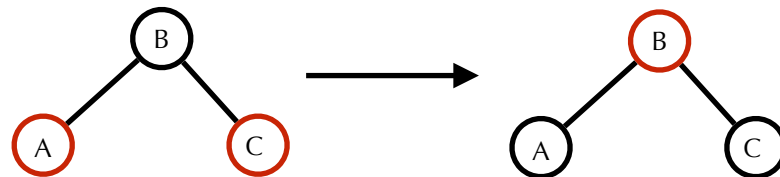
rotateleft



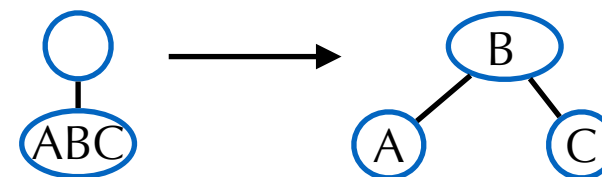
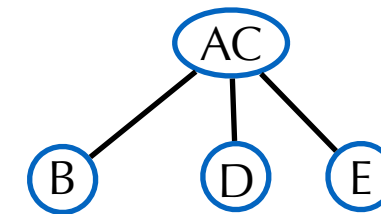
rotateright



flipcolors



## 2-3 tree



Check that  
order &  
black balance  
are preserved

Delete more  
complicated

# Worst-Case Analysis

**Prop.** The height of a red-black BST with  $n$  nodes is bounded by  $2 \log_2 n$ .

Proof: exercise.

## Summary

algorithm (data structure)	worst-case cost (after N inserts)		average-case cost (after N random inserts)	
	search	insert	search hit	insert
<i>sequential search (unordered linked list)</i>	$N$	$N$	$N/2$	$N$
<i>binary search (ordered array)</i>	$\lg N$	$N$	$\lg N$	$N/2$
<i>binary tree search (BST)</i>	$N$	$N$	$1.39 \lg N$	$1.39 \lg N$
<i>2-3 tree search (red-black BST)</i>	$2 \lg N$	$2 \lg N$	$1.00 \lg N$	$1.00 \lg N$

Empirical.  
No proof yet.



# References for this lecture

The slides are designed to be self-contained.

They were prepared using the following book that I recommend if you want to learn more:



# Next

Assignment: balanced BST

Next tutorial: Sudoku by WalkSat

Next week: String Algorithms 1

# Feedback

Moodle for the slides, tutorials and exercises.

Questions or comments: [Bruno.Salvy@inria.fr](mailto:Bruno.Salvy@inria.fr)