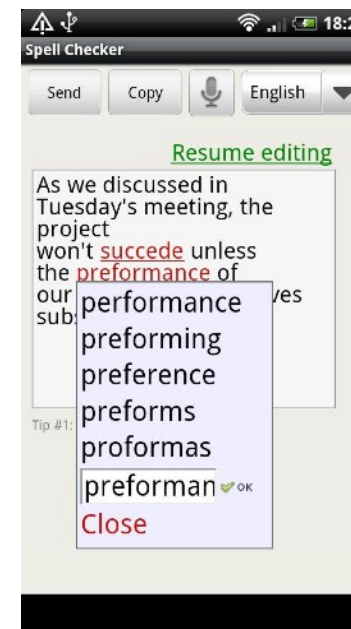
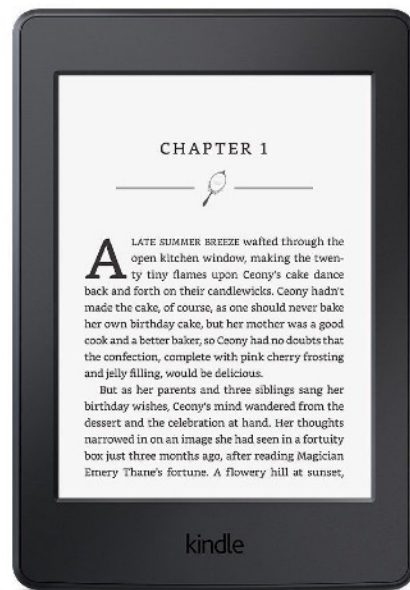


CSE202
Design and Analysis of Algorithms

Week 11 — String Algorithms 1 - Search

Strings Everywhere



```
class Node:
```

```
    def __init__(self, key, left=None, right=None):
        self.key = key
        self.left = left
        self.right = right
        self.size = 1
```

```
    def __str__(self):
        if self is None: return ''
        return str(self.left) + str(self.key) + str(self.right)
```

```
class BST:
```

```
    def __init__(self):
```

Definitions:

letter=character=symbol (usually 7,8, or 16 bits);

alphabet: set of letters (often denoted Σ and $R = |\Sigma|$);

string=word=text: *finite* sequence of letters;

length=size of a word: number of letters.

Substring Search

Notation:

$$|T| = n,$$
$$|P| = m.$$

Input: two strings (text T and pattern P)

Output: answer to “is P a substring of T ?”

$$\exists i, \forall j \in \{0, \dots, m-1\}, T_{i+j} = P_j.$$

Aim: *small number of character accesses*

Algorithms of the Day:

	worst case	average case
Brute Force	$\leq nm$	$\leq 2(n - m + 1)$
Knuth-Morris-Pratt	$\leq n + m$	$\geq n$
Boyer-Moore	$\leq 3n$	$\approx n/m$

difficult, not done here

Recall also Rabin-Karp from tutorial 7

I. Brute Force

Algorithm

```
def bruteforce(text, pattern):  
    for i in range(len(text)-len(pattern)):  
        for j in range(len(pattern)):  
            if text[i+j]!=pattern[j]: break  
        else: return i  
    return -1
```

Worst-case: $P = a^{m-1}b, T = a^{n-1}b$

$\rightarrow m(n - m + 1)$ comparisons

Expected Number of Comparisons for All Matches

Fixed pattern, uniform random text

texts of length n having at least the k first letters of the pattern at a given location: R^{n-k}

having exactly the k first letters: $R^{n-k} - R^{n-k-1}$

comparisons at this location: $\sum_{k=0}^m (k+1)(R^{n-k} - R^{n-k-1})$
 $\leq \frac{R^n}{1 - 1/R}$

comparisons at all locations: $\leq \frac{(n - m + 1)R^n}{1 - 1/R}$

Expectation $\leq \frac{n - m + 1}{1 - 1/R}$

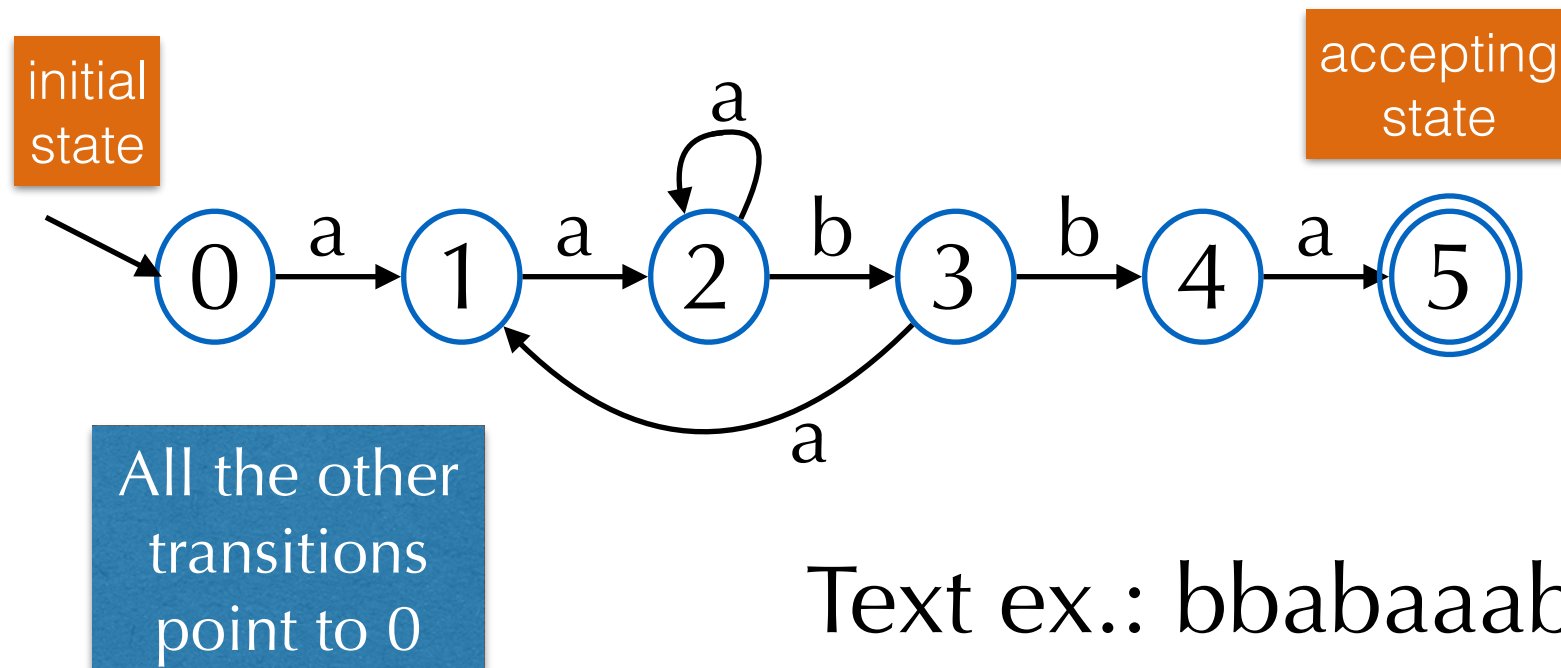
some texts
are counted
several times

II. Knuth-Morris-Pratt

Exploit the structure of the pattern

Pattern Compiled into Automaton

Automaton for aabba



```
def kmp(text,dfa):  
    m=len(dfa)  
    s=0  
    for i in range(len(text)):  
        s=dfa[s].get(text[i],0)  
        if s==m: return i  
    return -1
```

If the alphabet is small,
dictionaries are
replaced by arrays.


Def. Deterministic Finite Automaton

- a finite set Q of *states*;
- a *transition function* $\delta : Q \times \Sigma \rightarrow Q$;
- an *initial* state;
- one or several *accepting* states.

Each letter of the text is
accessed once

KMP works on streams:
never looks back

Computation of the Transitions



When a match fails at index i ,
 $i - 1$ characters of the text are known
→ imagine starting over from the 2nd one

```
def preprocess(pattern):  
    m=len(pattern)  
    dfa=[{} for i in range(m)]  
    dfa[0][pattern[0]]=1  
    state=0  
    for i in range(1,m):  
        for key in dfa[state]: dfa[i][key]=dfa[state][key]  
        state=dfa[state].get(pattern[i],0)  
        dfa[i][pattern[i]]=i+1  
    return dfa
```

state reads
pattern[1:m]

$O(m)$ operations

Exercise:
execute it step-by-step on this example

III. Boyer-Moore

The Boyer–Moore algorithm is considered as the most efficient string matching algorithm in usual applications. Crochemore Hancart Lecroq (2003)

Idea: Read from the End

Text: To be, or not to be, ...
Pattern: not to not to not to

last character shift

Boyer-Moore shift (defined later)

```
def bm(text, pattern, lcs, bms):  
    n = len(text)  
    m = len(pattern)  
    i = 0  
    while i <= n - m:  
        for j in range(m - 1, -1, -1):  
            if text[i + j] != pattern[j]:  
                i += max(1, j - lcs.get(text[i + j], -1), bms[j])  
                break  
            else: return i  
    return -1
```

```
def lastoccurrence(pattern):  
    m = len(pattern)  
    lcs = {}  
    for i in range(m - 1):  
        lcs[pattern[i]] = i  
    return lcs
```

Worst-case for last character heuristic:

$$P = ba^{m-1}, T = a^n$$

Worst-case becomes linear
with Boyer-Moore shift

Average-Case Complexity of the Last Character Heuristic

Fixed pattern, uniform random text, $m \leq R$

$$\mathbb{E}(\text{\#comparisons}) \approx \frac{n}{m} \text{ for large } R/m$$

n/m is
optimal

Also observed in practice

Assignment: step-by-step proof

Shift by Longest Suffixes

Text:

abaab**b**baa**b**ab**b**aaabaabaab**b**abaabaa

Pattern:

abaababaabaa

abaababaabaa

abaababaabaa

abaababaabaa

abaababaabaa

Find the smallest shift compatible with the letters read so far

$$\text{bms}[j] := \min \left\{ s > 0 \mid \left(\forall k \in \{j+1, \dots, m\}, s > k \text{ or } P[k-s] = P[k] \right) \right. \\ \left. \text{and } (s > j \text{ or } P[j-s] \neq P[j]) \right\}$$

Pattern	a	b	a	a	b	a	b	a	a	b	a	a
Longest suffix	1	0	1	4	0	1	0	1	4	0	1	
BM shift	8	8	8	8	8	8	8	3	11	11	1	2

slide the pattern to the left over itself and measure overlap

a. prefixes that are also suffixes
b. internal overlap

2. shift leftmost abaa — next abaa — 1. leftmost a

Corresponding Code

Linear
complexity
possible,
but harder

```
def longestsuffix(pattern):  
    m=len(pattern)  
    ls=[0]*(m-1)  
    for i in range(m-2,-1,-1):  
        for j in range(i+1):  
            if pattern[m-1-j]==pattern[i-j]:  
                ls[i] += 1  
            else: break  
    return ls
```


Linear
complexity

```
def bmshift(pattern):  
    ls=longestsuffix(pattern)  
    m=len(pattern)  
    bms=[m]*m          # init: default shift m  
    j=0  
    for i in range(m-2,-1,-1):  
        if ls[i]==i+1:  # a prefix is a suffix  
            for j in range(j,m-i-1): bms[j]=m-1-i  
    for i in range(m-1): # rightmost match  
        bms[m-1-ls[i]]=m-1-i  
    return bms
```

References for this lecture

The slides are designed to be self-contained.

They were prepared using the following books that I recommend if you want to learn more:



Next

Assignment: small alphabets

Next tutorial: searching for regular expressions

Next week: String Algorithms 2 — Compression

Feedback

Moodle for the slides, tutorials and exercises.

Questions or comments: Bruno.Salvy@inria.fr