

Branch-and-bound for TSP

In this tutorial we consider the TSP problem (traveling salesman, we formulate here the problem on weighted graphs rather than on matrices). A weighted graph G is a graph where each edge carries a weight in \mathbb{R} . A Hamiltonian circuit of G is a cycle of edges of G that visits each vertex of G exactly once. The cost of such a circuit is the sum of weights of its edges. The TSP problem is the optimization problem that consists in finding a Hamiltonian circuit of minimal cost (if no Hamiltonian circuit exists, the answer is $+\infty$). For example Figure 1 shows¹ a weighted graph on the left side, and the optimal Hamiltonian circuit (which has cost 8) on the right side.




FIGURE 1 – Left : a weighted graph. Right : the optimal Hamiltonian circuit, which has cost 8.

Download the files `WG.py`, `UF.py`, `test.py` and `cities.txt` (the file `cities.txt` contains a matrix of distances between european cities, to be used by a test later on, the file `UF.py` just contains a UnionFind class that is used by the minimal spanning tree method in `WG`, all the code to be completed is in `WG.py`). The file `WG.py` contains a class where each instance represents a weighted graph. It has a constructor that receives a list such as `L=[['a','b',1],['c','b',2],['a','c',4]]`, which is the list of edges given with their weights (for instance `a` and `c` are connected by an edge of weight 4). The list is stored in the attribute `self.edges` (sorted by edge-weights, to facilitate the computation of minimum spanning trees). The constructor also builds an attribute `adj` which is a dictionary associating to vertices their neighbors and the weights of edges between them. In particular : testing if `x` is a vertex of `self` is done by `if x in self.adj`. If `x, y` are two vertices, testing if `x` and `y` are adjacent is done by `if y in self.adj[x]`, and if true, `self.adj[x][y]` gives the weight of the edge connecting `x` and `y`.

1 A brute force recursive algorithm

Question 1. The method `min_cycle_aux(self,w,L,S)` has the following parameters : `L` is a list of vertices (of length at least 1) that forms a path in the graph, `w` is the total weight of edges on this path, and `S` is the set (Python `set` structure) of vertices that are not in `L`. The method has to return a pair `(W,Cyc)` where `Cyc` is a Hamiltonian circuit of smallest cost among the Hamiltonian circuits that start with the path `L`; and `W` is the cost of `Cyc`. For instance, in the example of Figure 1, if we take `L=[A,F,G]`, then there are two Hamiltonian circuits that extend `L` : the circuit `[A,F,G,D,E,B,C,H,A]` of cost 14, and the circuit `[A,F,G,H,C,D,E,B,A]` of cost 15. Hence the method has to return the pair `(14,[A,F,G,D,E,B,C,H,A])`. Complete the method `min_cycle_aux(self,w,L,S)` (which has to proceed recursively). To test your method execute `test1()`.

1. The two figures are taken from Section 9.1.2 of the book “Algorithms” by Dasgupta, Papadimitriou, Vazirani, which we closely follow.

Question 2. Then complete the method `min_cycle(self)` that has to return a pair (W, Cyc) where Cyc is a Hamiltonian circuit of smallest cost, and W is its cost (it can help to write an auxiliary method that returns the set of vertices of the graph). To test your method execute `test2()`.

2 Improvement using branch-and-bound




FIGURE 2 – The exploration tree for the graph of Figure 1 (starting from vertex A). By each node, corresponding to a path L starting from A , the framed number indicates $\text{low}(L)$. The tree hanging from the node is not explored if $\text{low}(L)$ is at least as large as the minimal cost of a Hamiltonian circuit seen so far. This value called *bestsofar*, is equal to $+\infty$ up to reaching the node for $L = [A, B, C, D, E, F, G, H]$ where it becomes equal to 11, and it stays equal to 11 up to reaching the node for $L = [A, F, E, B, C, D, G, H]$ where it becomes 8 (the tree is explored according to a depth-first search from left to right).

Question 3 (Exercise). With the notation of Question 1, and assuming S is not empty, let T_S be a spanning tree of minimal cost among the trees whose spanned vertices are those of S , and let w_S be the cost of T_S ($w_S = +\infty$ if there is no such tree, i.e., if the graph restricted to S is not connected). Let w_{start} be the minimal weight over edges from $L[0]$ to vertices in S ($w_{\text{start}} = +\infty$ if there is no such edge), and let w_{end} be the minimal weight over edges from $L[-1]$ to vertices in S ($w_{\text{end}} = +\infty$ if there is no such edge). Show that the weight W of the output circuit Cyc satisfies

$$W \geq w + w_{\text{start}} + w_{\text{end}} + w_S.$$

(hint : note that a path connecting all vertices of S is a special kind of spanning tree on S)

In other words, if we let $\text{low}(L)$ be the quantity $w + w_{\text{start}} + w_{\text{end}} + w_S$, then $\text{low}(L)$ provides a *lower bound* on the value of any circuit that extends L , and moreover this lower bound can be computed fast (we use here Kruskal's algorithm to compute minimal spanning trees). The recursive method in Question 1 can be represented as the exploration of a tree : starting the exploration from a vertex (vertex A for instance), each call to `min_cycle_aux` corresponds to a node of the tree, and the path L from the root to the node corresponds to the list L that is the parameter of `min_cycle_aux` for that call. A crucial observation is that, if during the exploration we are at a node L such that $\text{low}(L)$ is at

least as large as the mincost over Hamiltonian circuits seen so far, then we can abort the exploration at that node (i.e., avoid exploring the tree hanging from L). This is illustrated in Figure 2.

Question 4. Complete the method `lower_bound(self,w,L,S)` that computes $\text{low}(L)$ (with w the cost of L , and S the set of vertices that are not in L , we assume S is not empty). You will have to use the method `weight_min_tree(self,S)` in the class `WG` (this method returns the quantity w_S). To test your method execute `test4()`.

Question 5. Complete the method `min_cycle_aux_using_bound(self,bestsofar,w,L,S)`. It has to return `(math.inf,[])` if $\text{low}(L) \geq \text{bestsofar}$, and otherwise it has to return the same output as the method `min_cycle_aux(self,w,L,S)`, while implementing the branch-and-bound technique illustrated in Figure 2. We assume that the method will be called with the value of `bestsofar` giving the mincost over Hamiltonian circuits seen so far (i.e., over the nodes visited before L), accordingly you have to correctly update the value of `bestsofar` for the recursive calls to `min_cycle_aux_using_bound` in your code. Then complete the method `min_cycle_using_bound(self)` that has to return the same output as `min_cycle(self)`.

To test your method, you can use two methods : `test_random_graph(n)` that tests both methods `min_cycle` and `min_cycle_using_bound` (and compares their running times), for a complete graph on n vertices where the edge-weights are (uniformly) random in $[0, 1]$. The other method is `test_trip()` that runs `min_cycle_using_bound` to determine the shortest circuit to visit a selected set of cities in Europe (among a fixed list of cities).

Remark. To evaluate the efficiency of the branch-and-bound technique, you can also introduce a global variable `nr_calls` and use it to count the total number of (recursive) calls to `min_cycle_aux` during the execution of `min_cycle`. Doing the same for the methods using bounds, you can then compare the number of calls in `min_cycle` and in `min_cycle_using_bound` for random graphs with n vertices.

Remark. The efficiency of the branch-and-bound technique crucially depends on the quality of the lower bounds, and there is a delicate trade-off (finding a lower bound that is good and not too time-consuming to compute). In our case, we can generalize the lower bounds based on the following simple observation due to Held and Karp : if we fix any function $\pi : V \rightarrow \mathbb{R}$ (with V the vertex-set) and update the weight of each edge $e = (u, v)$ to be $w_\pi(e) = w(e) + \pi(u) + \pi(v)$, then the optimal values are clearly related by $\text{TSP} = \text{TSP}_\pi - 2 \sum_{v \in V} \pi(v)$. If we are at a certain node of the exploration tree we can use this observation and look for functions π that will give us better and better lower bounds (note that the optimal spanning tree on vertices of S will depend on the function π that is chosen). This can be efficiently done by an iteration updating π (improving the lower bound) progressively, based on a gradient computation.

Several other clever techniques have been developed to gain on time-complexity (looking for local optima under certain local operations to update the circuit, Linear Programming relaxation combined with branch-and-bound techniques, etc.), see for instance Section 9.10.2 of the book ‘The nature of computation’ by Moore & Mertens for a clear presentation of such techniques. This allows to solve the TSP problem on some very large graphs, e.g. computing the shortest tour through all 24 978 cities of Sweden.